PRACTICAL PLAYBOOK

AI Adoption for California Cities

Six steps to adopt AI safely, smartly, and without the two-year detour.

Built from the guidance and real-world learnings of some of the most innovative jurisdictions in California.

Madison AI

www.madisonai.com

The 6-Step Process

Get Aligned	Run a working session with your team	90 minutes
Set Guardrails	Build your governance plan: principles, acceptable use, accountability	1 day
Start With Your Data	Identify your most authoritative records	1-2 weeks
Pick Your First Win	Choose one high-impact, high-pain use case	1 week
Train by Doing	Run a real task with your champion user	Ongoing
Expand When Ready	Add data, add use cases, add users	90+ days

STEP 1

Get Aligned

90 minutes. One room. Walk out with your top use cases and a reason to start.



“We don’t have enough people, and I don’t know any government agency that has extra people. The demands are just going up and up and up. We’re looking for ways to add onto the people we have and remove some of the drudgery work.”

Chris Gray, Deputy Executive Director, WRCOG

Your Mission Comes First. The Technology Comes Second.

AI does not start with software. It starts with your "why." Why does your organization exist? What impact are you trying to have on your community? AI should multiply that mission, not create a side project.

You do not need a 50-page strategy. You need one working session with your team to get aligned on three things: your direction, your biggest pain points, and your willingness to try something new.

Who Should Be in the Room

Your leadership team AND a handful of frontline staff. Management analysts. Clerks. Planners. Admin staff. The people closest to the problems.

AI transformation falls flat when it starts at the top and gets pushed down. Start from the bottom up. The people doing the work know which tasks eat their days.

1 The 5-Hour Question

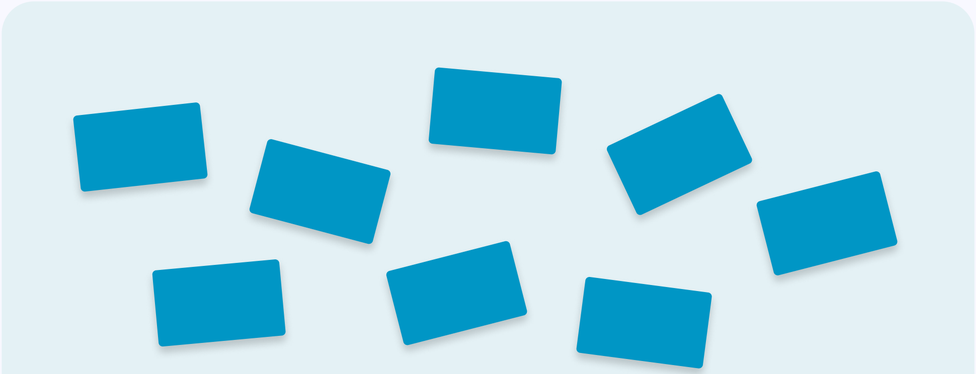
20 min | Sticky notes

Hand everyone a sticky note or index card. Ask one question:

"What are 3 things in your day-to-day role you could do differently if you had 5 extra hours a week?"

That is not hypothetical. Jurisdictions across California are getting exactly that kind of time back right now. So the question is not "can AI help?" It is "where would the time go?"

Give people **5 minutes** to write silently. Then go around the room. Collect every answer on a whiteboard or shared doc. No filtering. No debating. Just capture.



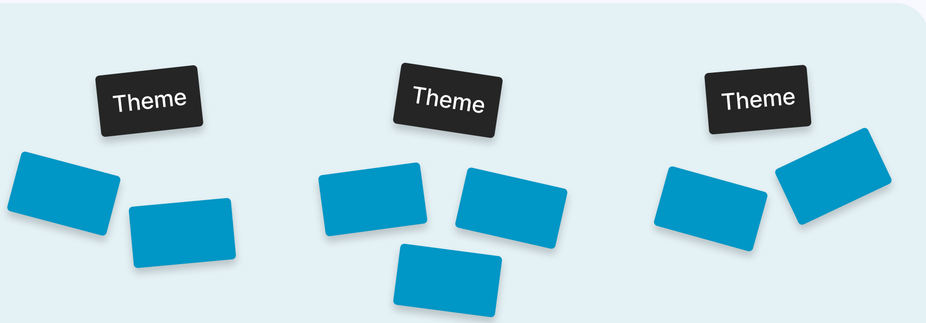
2 Find the Patterns That Matter

20 min | Whiteboard

Look at what the team wrote and cluster the answers.
You are looking for three types of tasks:

- **Repetitive tasks.** Things done the same way over and over. Formatting staff reports. Compiling prior actions. Pulling meeting histories.
- **Time sinks with low value.** Hours of research for something that should take minutes. Finding a past council vote. Tracking down a contract expiration date. Answering "when did we last...?"
- **Interrupt magnets.** Questions that pull colleagues away from their actual jobs because the answer lives in someone's head or is buried in 14 different file folders.

Circle the clusters. These are your candidate use cases.



3 Effort & Impact Grid

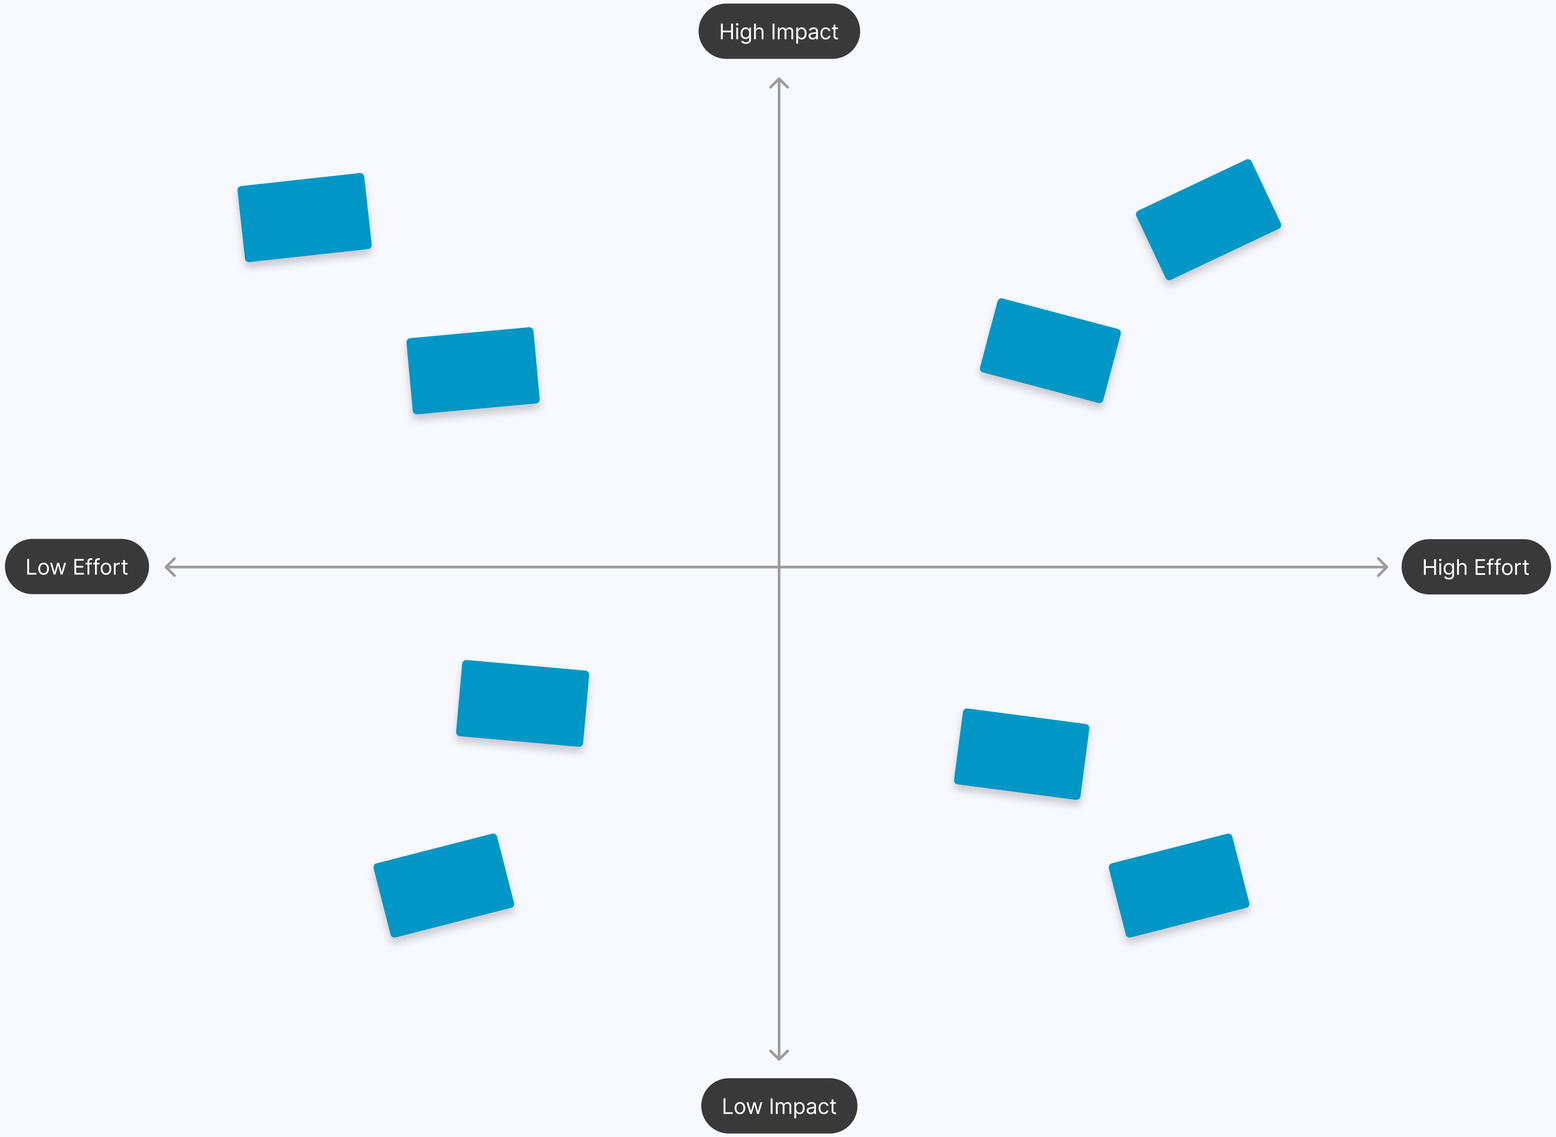
20 min | Whiteboard

On a whiteboard, draw the grid shown on the right. Take your candidate use cases from Part 2 and place each one on the grid.

For each one, ask the room two questions:

"Is this low or high effort to implement?"
"Is the impact on our team low or high?"

Do not overthink it. Gut calls are fine.
You are not committing to anything.
You are prioritizing.



4 Make the Case in One Sentence

15 min | Pen & paper

For each opportunity in the upper left quadrant, write a one-sentence case using one of these frames:

"Improving the current:" How does this make an existing process faster, cheaper, or more reliable?

"Creating the future:" How does this help us serve our community in ways we could not before?

Most first use cases will be "improving the current." That is exactly right. Start there. The transformation opportunities show up naturally once the system is running and people start asking "what else can it do?"

5 Plant the Bigger Seeds

15 min | Optional

Two questions to leave on the whiteboard for later:

- **Is there a process that is ripe for reinvention, not just acceleration?**
- **Could faster internal research lead to better community outcomes, like faster permit turnarounds or smoother staff transitions?**

You do not need to answer these today.
Writing them down keeps the door open.

What you need to do:

- Block 90 minutes this month.
  Run Parts 1 through 4 with your team.
- Walk out with your top 2 or 3 low-effort, high-impact use cases and a one-sentence justification for each.

That is your starting point.

STEP 2

Set Guardrails

Every government needs an AI governance plan.
It does not have to be 100 pages. But you need one.



“For a small staff, AI lets us be nimble: adopt a policy, train the team, ingest the data, and suddenly everyone can find the answers.”

Dr. Peter Pirnejad, City Manager (Ret.), Town of Los Altos Hills

Why This Is Not Optional

Here is the hard truth: without an AI governance plan, you will not realize the benefits AI can bring your organization. You will constantly hit the same five walls: values misalignment, data exposure, bias, staff afraid to experiment, and staff afraid they will lose their jobs.

An AI governance plan is not a bureaucratic exercise. It is the thing that makes everything else possible. It creates the safety for your team to actually use the tools. Los Altos Hills did it with a 7-page policy. Took it to committee, took it to council, had it adopted, and started training staff. You can do this.

Your Governance Plan Has Three Parts:

1. Guiding Principles and Ethics.
2. What AI Is For. What It Is Not For.
3. How You Create Accountability

1 Guiding Principles and Ethics

Your guiding principles are **3 to 5 statements** that tell every employee what responsible AI use looks like. Anchor them to your existing core values. Write them so clearly that anyone can read them and know whether they are in compliance.

Five themes your principles should address:

Data protection.

How is data used? What data can and cannot be exposed to AI tools? If you are not clear here, someone will put sensitive information into an open-source tool on day one.

Bias prevention.

AI can replicate and amplify biases already present in your data. State how your organization will identify and prevent that. Human oversight is the key.

Job security.

This is the one that matters most for adoption. If your staff thinks AI will replace them, they will not participate. Be explicit: AI gives your team superpowers. It does not eliminate roles.

Space for experimentation.

Your team needs permission to try, fail, and learn. The guardrails exist so people can move fast inside them.

Original content and IP.

If your team uses AI to draft documents, how much needs to be original? What is acceptable use versus plagiarism? Get clear on this early.

How to build them

- Anchor to your core values
- Identify which themes matter most
- Write 3 to 5 clear statements in the positive ("We will...")

For a detailed walkthrough, see:



How to Develop AI Guiding Principles

Common Pitfalls

Don’t suck in all of SharePoint
Start with adopted records, not draft folders. The AI treats every document as authoritative. Curate first, ingest second.

Don’t skip governance to move fast
A one-page set of guiding principles takes a day. Rebuilding trust after a bad output reaches the public takes much longer.

Don’t expect perfection from day one
Prompts need refining. Data conversion takes iteration. The best results come from continuous improvement, not one-time setup.

Don’t plan top-down only
Run a short workshop: ask each person what 3 things they’d do differently with 5 extra hours a week. That’s where the gold is.

Don’t let IT spec your AI alone
AI strategy is one part technology and four parts people. Department heads and frontline staff need to be at the table. The best use case ideas come from the people doing the work.

Don’t leave jobs ambiguous
If your team thinks AI might replace them, they won’t participate. Be explicit in your guiding principles: AI enhances, it doesn’t eliminate roles.

Don’t confuse outputs with answers
AI generates responses. Those responses are often excellent, but they’re not gospel. The better the AI gets, the harder the remaining errors are to spot. Verification standards on day 100 should be as high as day 1.

The Bottom Line

Stop overthinking it. The jurisdictions seeing results didn’t wait for perfect conditions. **They picked a pain point, used clean data, set basic guardrails, and started.**

Every single one says the same thing: they wish they’d started sooner.

You don’t need a massive budget. You don’t need a year-long procurement. You don’t need a dedicated IT team. **You need one hour of leadership alignment, one day of governance, one clean data source, and one use case that matters.** The rest builds from there.

If you’re hesitant, your staff will be hesitant. If your staff is hesitant, they’re falling behind.

And the institutional knowledge walking out your door isn’t waiting for you to be ready.

2 What AI Is For. What It Is Not For.

Draw the line clearly so your team knows where to run.

Use AI for:

- research and search across institutional records
- drafting documents in your format and style
- summarizing meeting minutes and board actions
- answering factual questions grounded in your adopted data
- accelerating repetitive administrative tasks

Do not use AI for:

- making final decisions without human review
- processing confidential personnel or legal information through unsecured tools
- generating public-facing content without staff verification
- replacing the critical thinking and judgment your staff brings to every decision

A note on ChatGPT and open-source tools

Many jurisdictions draw a clear line between purpose-built AI (trained on your data, secured in your environment) and general-purpose tools like ChatGPT (trained on the open internet).

The City of San Jose's AI Policy explicitly requires staff verification of all AI content before public release and tracks how staff use AI tools. Your policy should address both types of tools and set clear boundaries.

3 How You Create Accountability

Someone has to own this. Name them.

The rule is simple: staff own every output. AI is a tool, like a spreadsheet. If the output is wrong, that is on the person, not the machine. Chris Gray at WRCOG put it plainly: "I don't care if you write a staff report yourself or get it from the AI. You are responsible for the content. Period."

> “I don’t care if you write a staff report yourself or get it from the AI. You are responsible for the content. You can’t turn around and blame it on the AI.”
>
> Chris Gray, Deputy Executive Director, WRCOG

But personal accountability is not enough on its own. You also need an organizational structure. Someone needs to oversee how AI is used, make sure the guardrails hold, and evolve the policy as the technology changes. Governance without names attached to it is just a document nobody enforces.

Pick a governance structure

There are three models. Pick the one that fits how your organization actually works.

Centralized

One person or team owns all AI oversight. Best for hierarchical organizations with tight data controls or heavy regulatory requirements. One executive champion, one set of rules for the whole organization.

Decentralized

Each department self-governs under shared guiding principles. Best for flat, nimble organizations with high staff engagement. Gives teams flexibility to adapt AI to their specific workflows while staying inside the guardrails.

Center of Excellence

A hybrid. Centralized AI expertise and best practices, with room for departments to innovate on their own. This is where most mid-sized cities and counties land. Avoids bottlenecking everything through one team while still keeping strategic alignment.

Whichever you choose, fill these roles: an executive champion who clears the path, an oversight lead who runs the process, a technical lead who manages the tools, and a legal lead who advises on risk. In a 30-person organization, one person might wear two or three of these hats. That is fine. Just make sure someone is named.

3 How You Create Accountability

Pick one enforcement mechanism

You have your principles. You have your structure. Now make accountability visible to every employee.

Require a policy acknowledgment before access

The City of San Jose requires staff to follow its official AI Policy (City Policy Manual 1.7.12) and tracks how staff use AI tools to improve training and safety. All AI-generated content must be checked by staff before being shared publicly.

Require training before access

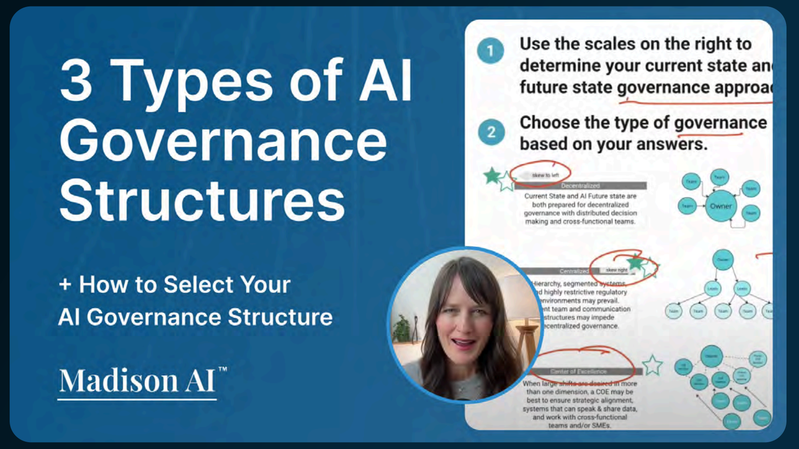
Alameda County's GenAI Policy requires all staff to complete AI training administered by their IT department before using generative AI tools. California's state-level guidelines similarly require agency executives, technical staff, and government workers to receive AI training covering best practices and discrimination prevention.

Add a disclaimer to AI-drafted documents

The City of Sunnyvale adds language at the top of AI-drafted staff reports stating that the author is responsible for verifying all content before it is published or submitted.

You do not need all three. Pick one. Get it in place.
You can layer on more as your team matures.

For a deeper dive on assessing your organization and selecting the right model:



How to Select Your AI Governance Structure

You Do Not Have to Build This From Scratch

The GovAI Coalition (led by the City of San Jose) is a network of 600+ public servants from 250+ governments. They have published free, open-source governance templates built by practitioners, for practitioners, aligned with NIST's AI Risk Management Framework.

Templates include: AI Use Policy, AI Policy Manual, AI Incident Response Plan, AI FactSheet templates for vendor transparency, vendor agreement templates, and use case templates with real examples. All free. All designed to be borrowed, edited, and adopted.

Resources from The GovAI Coalition

Madison AI's governance resources provide step-by-step guides for each piece of your plan:

- AI Governance Blueprint
- How to Develop AI Guiding Principles
- How to Select Your AI Governance Structure

Quick-Start Governance Checklist

- [ ] We have 3 to 5 AI guiding principles anchored to our core values.
- [ ] We have a clear position on data protection, bias, job security, experimentation, and IP.
- [ ] We have defined what AI should and should not be used for.
- [ ] We have picked a governance structure (centralized, decentralized, or center of excellence).
- [ ] We have named an executive champion, oversight lead, technical lead, and legal lead.
- [ ] We have one enforcement mechanism in place (policy acknowledgment, training, or disclaimer).
- [ ] Staff understand they own every output, no exceptions.
- [ ] Our policy has been reviewed by leadership and ideally adopted by council or board.

STEP 3

Start With Your Data

Your data is better than you think. The question is not "is it ready?" It is "where do we start?"



“Problem with local government is that we have data that’s complex, but it’s not big. We don’t have big data. We have complex, dirty data. And AI that can actually read and decipher PDFs and Laserfiche and videos — it was a game changer.”

Dr. Peter Pirnejad, City Manager (Ret.), Town of Los Altos Hills

The Myth of Dirty Government Data

There is a common misconception that government data is a mess. Too dirty, too unstructured, too scattered to work with AI. That you need a massive data cleanup project before you can even begin.

This is mostly wrong.

Yes, your data is unstructured. It lives in PDFs, scanned documents, Laserfiche, CivicClerk, agenda management systems, SharePoint, and filing cabinets. It is not in neat spreadsheets. But here is what most people miss: government data is actually some of the most complete, well-documented enterprise data that exists. You keep everything. You archive everything. Your council minutes, resolutions, staff reports, budgets, and ordinances go through a formal public process before they are adopted. That means they are reviewed, approved, and preserved as the official record.

That is not dirty data. That is a goldmine.

Generative AI was built for exactly this kind of unstructured information. Unlike older technologies that needed clean, structured databases to function, modern AI can read and understand PDFs, scanned documents, meeting minutes, and even handwritten notes. The technology has caught up to the way government actually stores information.

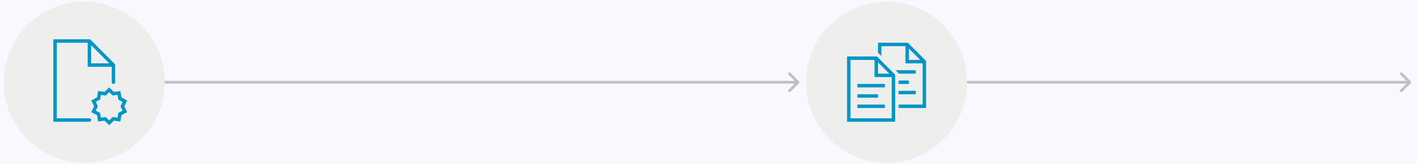
Your Data Is More Ready Than You Think

Before generative AI, working with unstructured data was nearly impossible at scale. You were limited to keyword searches (if you were lucky) and manual file digging. That is why institutional knowledge ended up trapped in people's heads rather than in any searchable system.

Now, AI can read across all of it. Pull it together. Cite where the answers came from. The barrier to entry is not a data cleanup project. It is simply choosing where to point the AI first.

Pick Your Starting Point

You do not need to load everything on day one.
You need one authoritative data source to begin with.



Start with the record of decision

Your adopted council or board records are almost always the best starting point. Minutes, resolutions, adopted staff reports, ordinances, and the attachments that went through the public process. This data has been reviewed, approved, and archived. It IS the institutional record. And it is usually already digital.

Expand from there

Once your core governance records are in, layer on additional sources: planning and zoning records, contracts and procurement files, HR policies, budget documents, community development files. Each new data source makes the AI smarter and more useful.

Common Starting Points by Department

City/County Clerk
Council minutes, resolutions, ordinances, adopted staff reports

Finance
Adopted budgets, CAFRs, audit reports

Planning
General plans, specific plans, zoning code, commission minutes

Procurement
Executed contracts, vendor agreements, RFPs

HR
Adopted personnel policies, MOUs

The One Thing to Get Right: Use Final Versions

The only place data quality really matters on day one is version control. You want to load the final, adopted versions of documents, not drafts.

One assistant county manager spent hours downloading files from SharePoint and got exactly what you would expect: the AI treated 14 drafts of the same document as 14 separate authoritative sources. The result was unreliable. Start with what your board or council has adopted and you avoid this problem entirely.

WRCOG learned this the useful way. Their AI surfaced the fact that staff had been referencing a July draft of their personnel manual instead of the August final that the board actually adopted. Nobody knew they were using the wrong version until the AI found the right one. That single discovery paid for the investment.

The rule: if it has been adopted by your governing body, load it. If it is a draft sitting in someone's shared drive, leave it out. You can always add more data later.

> “Nobody remembered we had adopted that resolution, even though everyone who was involved in passing it was still on staff. We found it in seconds. By the way, our attorney charges us a lot per hour — having him spend four or five hours looking for something is very expensive.”
>
> Chris Gray, Deputy Executive Director, WRCOG

Quick-Start Data Checklist

- ◯ We have identified our primary authoritative data source (usually clerk records or agenda management system).
- ◯ We are starting with board-adopted, final-version documents.
- ◯ We are NOT loading draft folders, shared drives, or unvetted content.
- ◯ We know which systems hold our key records (clerk, finance, planning, procurement, HR).
- ◯ We have identified who curates or manages each data source.
- ◯ We understand the data does not need to be "clean" in the traditional sense. It needs to be authoritative.

STEP 4

Pick Your First Win

Don’t boil the ocean. Pick the use case that will make believers out of skeptics.

The best first use case has three characteristics:

- It drives leadership crazy today
- It involves searching through years of records
- It currently takes hours or days when it should take minutes.

Top picks from jurisdictions that have done this:

Staff report research and drafting

The universal time sink. About 80% of a staff report is boilerplate: prior actions, background, format, section headers, strategic plan alignment. Let AI handle that. You focus on the 20% that requires critical thinking — the substance, the political sensitivity, the story you’re telling the board. Washoe County saves hundreds of hours per month. Fresno cuts prep from 6+ hours to under an hour.

“When did we last…?” questions

When did we last update the ADU code? What was the vote on the mayoral rotation? These used to require interrupting three different people and digging through Laserfiche. Now it’s a 30-second search with citations to the source document.

New employee onboarding

New hires who use the AI show up to meetings fully prepared — with budget history, voting context, and project timelines — without bothering six colleagues. One city manager noticed a new employee seemed unusually well-prepared. The answer: “I just researched it in the AI.”

Forgotten governance documents

WRCOG found a resolution nobody remembered passing, even though everyone involved was still on staff. The AI found it in seconds. That single find saved hours of attorney time.

The “first draft” challenge

Some staff will say the AI draft takes more effort to edit than writing from scratch. This is common early on.

The fix: invest time upfront defining what a good staff report looks like — section lengths, summary constraints, acronym formatting — so the AI matches your standards consistently.

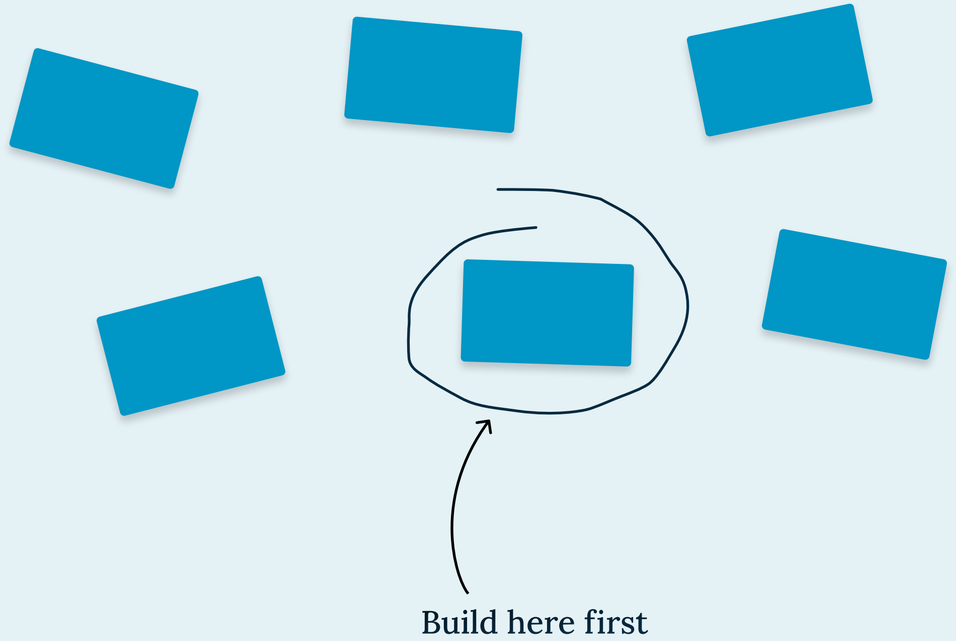


“About 80% of the content is boilerplate. Where we want people to spend most of their time is the real meat of the staff report — the discussion. The rest is just stuff you have to put in there.”

Chris Gray, Deputy Executive Director, WRCOG

ASK YOUR TEAM

“What’s the question you get asked most that takes the longest to answer?”



STEP 5

Train by Doing

Don’t over-train before launch.
Launch, then train through use.



“We’ve gone through a lot of iterations to get where we need to be. We see the value in it, but it has taken a lot of training and a lot of work.”

Dr. Peter Pirnejad, City Manager (Ret.), Town of Los Altos Hills

The most effective training isn’t a classroom session

It’s running a real task with a real user and walking through the results together. A few non-negotiables:

Always verify

The AI cites its sources — click through and confirm. Make it muscle memory. As one practitioner put it: “Start with AI, but end with the human eye.”

Ask it “why”

If you get an unexpected answer, ask the AI why it gave that response. It will explain which sources it used. One leader got two different answers from two slightly different prompts, then asked the AI to explain. It turned out one triggered a project-specific figure while the other pulled a program-level number. That moment became a training lesson for the whole team.

Staff own the output

No exceptions. AI is no different than a spreadsheet: if you type the wrong formula, you don’t blame Excel.

What to expect

- Some people will insist it's faster to do it themselves. They might be right — for now. That changes as data improves and prompting skills develop.
- AI gives different answers to slightly different prompts. That's not a bug. Train staff to be precise.
- Planners will want to know where every data point came from. Good — the citation links are there for that.
- Not every department should use every feature. Washoe County's Community Services skipped report drafting intentionally to preserve learning-by-doing. That's a valid choice.

The generational gap

Experienced leaders question AI outputs instinctively. Younger staff may accept them at face value.

Both tendencies need managing. Build a culture where verification is the norm regardless of age or comfort level.

What you need to do:

- Run your first real task with your champion user
- Walk through the results together

That's your training.

STEP 6

Expand When Ready

You started with one use case.
Now you know what works.

When you're ready, move on to your next priority project you identified earlier.

Build role-specific assistants

One for executives, one for planners, one for procurement, one for elected officials.

Solve temporary problems flexibly

PRA spike? Build a PRA assistant. Problem goes away? Turn it off. You don't need a permanent solution for a passing problem.

Capture institutional knowledge

Have departing leaders record short video journals about key decisions. Transcribe and ingest. That knowledge becomes permanently searchable.

Frame every expansion in two categories:

Here are some ideas for where to take things from here.

Improving the current

Making an existing process faster, cheaper, or more reliable.

Creating the future

Serving your community in a way you couldn't before.

Most jurisdictions start with "improving the current" and naturally discover "creating the future" opportunities along the way.

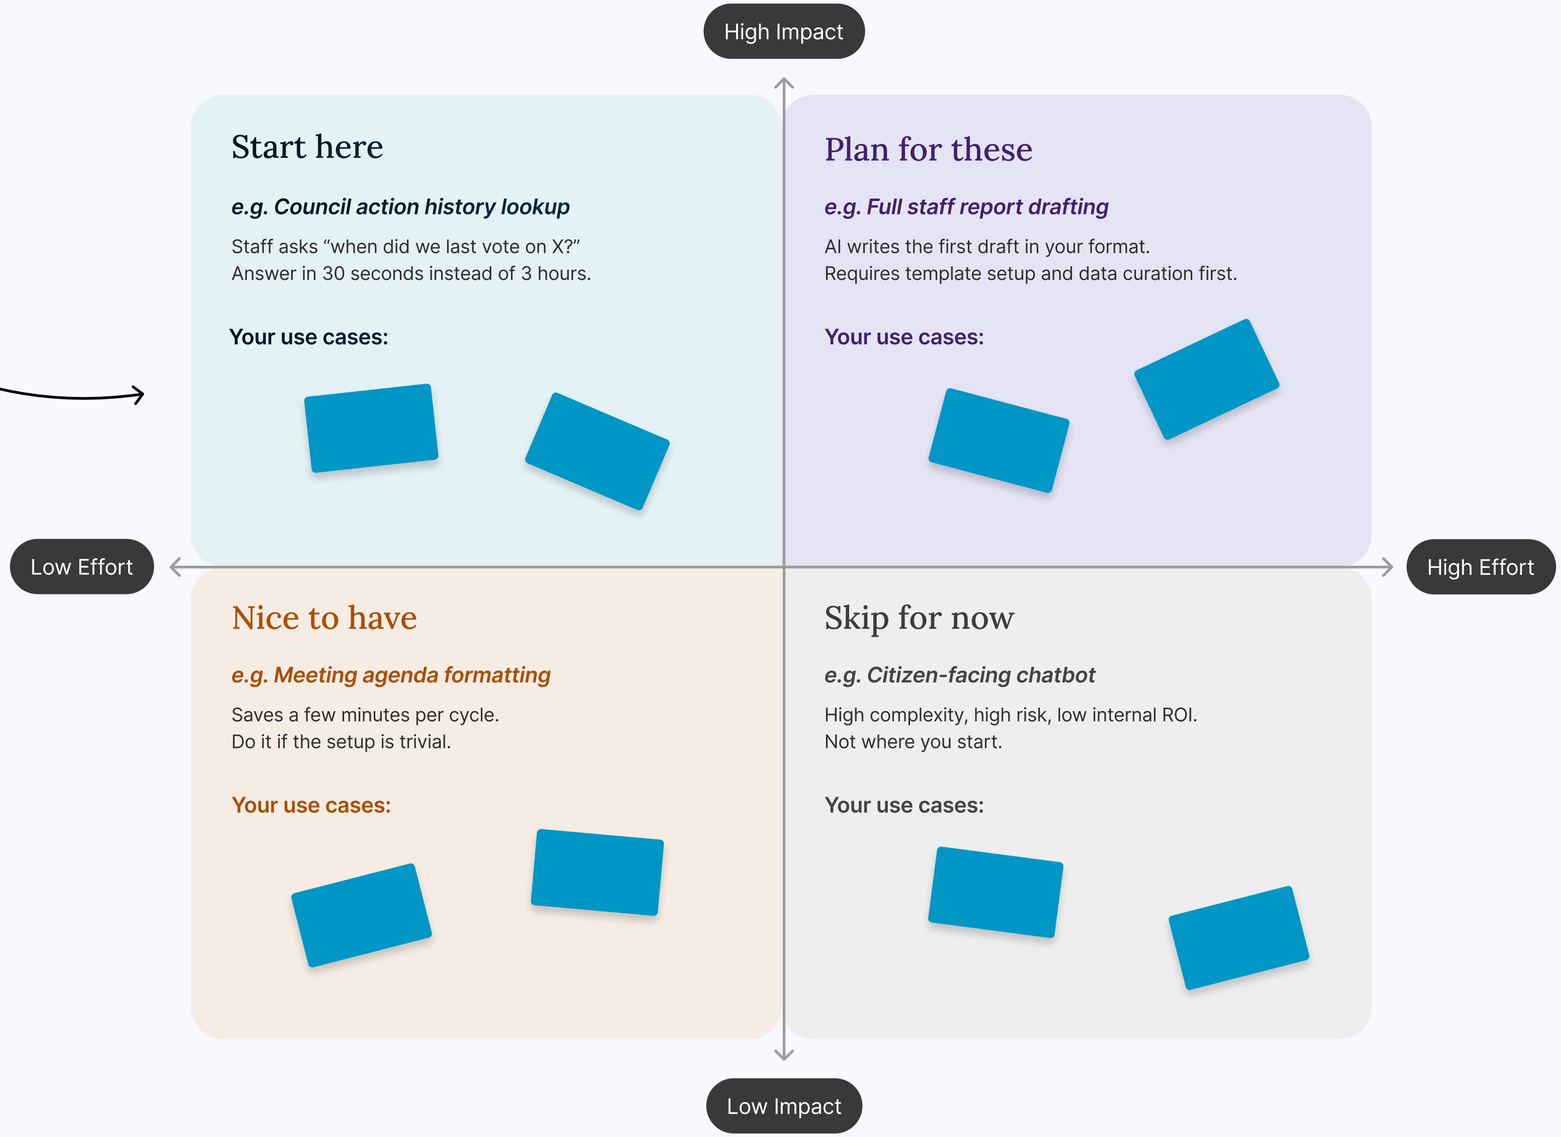
What you need to do:

After 90 days, ask your team:

"What else do you wish it could do?"

That's your roadmap.

The upper left quadrant is your starting point. Low effort, high impact. That is your low-hanging fruit. That is where you build your first win and make believers out of skeptics.



Madison AI

www.madisonai.com

Thank you!

Your Government's AI Built from Your Institutional Knowledge and Data.

Safe.

Secure.

Only your data.

Scan the QR code and book your demo today